

Artificial Intelligence Nanodegree

Convolutional Neural Networks

Project: Write an Algorithm for a Dog Identification App

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

Note: Once you have completed all of the code implementations, you need to finalize your work by exporting the iPython Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to '\n', "**File -> Download as -> HTML (.html)**". Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

Note: Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this IPython notebook.

Why We're Here

In this notebook, you will make the first steps towards developing an algorithm that could be used as part of a mobile or web app. At the end of this project, your code will accept any user-supplied image as input. If a dog is detected in the image, it will provide an estimate of the dog's breed. If a human is detected, it will provide an estimate of the dog breed that is most resembling. The image below displays potential sample output of your finished project (... but we expect that each student's algorithm will behave differently!).

In this real-world setting, you will need to piece together a series of models to perform different tasks; for instance, the algorithm that detects humans in an image will be different from the CNN that infers dog breed. There are many points of possible failure, and no perfect algorithm exists. Your imperfect solution will nonetheless create a fun user experience!

The Road Ahead

We break the notebook into separate steps. Feel free to use the links below to navigate the notebook.

- [Step 0](#): Import Datasets
 - [Step 1](#): Detect Humans
 - [Step 2](#): Detect Dogs
 - [Step 3](#): Create a CNN to Classify Dog Breeds (from Scratch)
 - [Step 4](#): Use a CNN to Classify Dog Breeds (using Transfer Learning)
 - [Step 5](#): Create a CNN to Classify Dog Breeds (using Transfer Learning)
 - [Step 6](#): Write your Algorithm
 - [Step 7](#): Test Your Algorithm
-

Step 0: Import Datasets

Import Dog Dataset

In the code cell below, we import a dataset of dog images. We populate a few variables through the use of the `load_files` function from the scikit-learn library:

- `train_files`, `valid_files`, `test_files` - numpy arrays containing file paths to images
- `train_targets`, `valid_targets`, `test_targets` - numpy arrays containing onehot-encoded classification labels
- `dog_names` - list of string-valued dog breed names for translating labels

In [88]:

```
from sklearn.datasets import load_files
from keras.utils import np_utils
import numpy as np
from glob import glob

# define function to load train, test, and validation datasets
def load_dataset(path):
    data = load_files(path)
    dog_files = np.array(data['filenames'])
    dog_targets = np_utils.to_categorical(np.array(data['target']), 133)
    return dog_files, dog_targets

# load train, test, and validation datasets
train_files, train_targets = load_dataset('dogImages/train')
valid_files, valid_targets = load_dataset('dogImages/valid')
test_files, test_targets = load_dataset('dogImages/test')

# load list of dog names
dog_names = [item[20:-1] for item in sorted(glob("dogImages/train/*/"))]
```

```
# print statistics about the dataset
print('There are %d total dog categories.' % len(dog_names))
print('There are %s total dog images.\n' % len(np.hstack([train_files, valid_files, test_files])))
print('There are %d training dog images.' % len(train_files))
print('There are %d validation dog images.' % len(valid_files))
print('There are %d test dog images.' % len(test_files))
```

There are 133 total dog categories.

There are 8351 total dog images.

There are 6680 training dog images.

There are 835 validation dog images.

There are 836 test dog images.

Import Human Dataset

In the code cell below, we import a dataset of human images, where the file paths are stored in the numpy array `human_files`.

In [89]:

```
import random
random.seed(8675309)

# load filenames in shuffled human dataset
human_files = np.array(glob("lfw/*/.*"))
random.shuffle(human_files)

# print statistics about the dataset
print('There are %d total human images.' % len(human_files))
```

There are 13233 total human images.

Step 1: Detect Humans

We use OpenCV's implementation of [Haar feature-based cascade classifiers](#) to detect human faces in images. OpenCV provides many pre-trained face detectors, stored as XML files on [github](#). We have downloaded one of these detectors and stored it in the `haarcascades` directory.

In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

In [90]:

```
import cv2
import matplotlib.pyplot as plt
%matplotlib inline

# extract pre-trained face detector
face_cascade =
cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')

# load color (BGR) image
img = cv2.imread(human_files[3])
```

```

# convert BGR image to grayscale
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)

# find faces in image
faces = face_cascade.detectMultiScale(gray)

# print number of faces detected in the image
print('Number of faces detected:', len(faces))


# get bounding box for each detected face
for (x,y,w,h) in faces:
    # add bounding box to color image
    cv2.rectangle(img, (x,y), (x+w,y+h), (255,0,0), 2)

# convert BGR image to RGB for plotting
cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

# display the image, along with bounding box
plt.imshow(cv_rgb)
plt.show()

```

Number of faces detected: 3



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

In [91]:

```

# returns "True" if face is detected in image stored at img path

```

```
def face_detector(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade.detectMultiScale(gray)
    return len(faces) > 0
```

(IMPLEMENTATION) Assess the Human Face Detector

Question 1: Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

Answer: There is 99 of human faces detected in the `human_files` and there is 11 of human faces detected in the `dog_files`. The results are really good for the human class, however there is also a significantly high detection rate for human faces in dog images. The method used here is based on the OPENCV Haar file: `haarcascade_frontalface_alt`.

In [92]:

```
human_files_short = human_files[:100]
dog_files_short = train_files[:100]
# Do NOT modify the code above this line.

## TODO: Test the performance of the face_detector algorithm
## on the images in human_files_short and dog_files_short.

countHumans = 0
countDogs = 0
for i in range(100):
    if face_detector(human_files_short[i]):
        countHumans += 1
    if face_detector(dog_files_short[i]):
        countDogs += 1

print('Alt1 Method: Number of faces in human 100 images: {}'.format(countHumans))
print('Alt1 Method: Number of faces in dog 100 images: {}'.format(countDogs))
```

Alt1 Method: Number of faces in human 100 images: 99

Alt1 Method: Number of faces in dog 100 images: 11

Question 2: This algorithmic choice necessitates that we communicate to the user that we accept human images only when they provide a clear view of a face (otherwise, we risk having unnecessarily frustrated users!). In your opinion, is this a reasonable expectation to pose on the user? If not, can you think of a way to detect humans in images that does not necessitate an image with a clearly presented face?

We suggest the face detector from OpenCV as a potential way to detect human images in your

algorithm, but you are free to explore other approaches, especially approaches that make use of deep

algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on each of the datasets.

Answer: Three additional methods were used to benchmark different approaches of human faces detections. All three methods are based on the OPENCV Haar library files and the results can be seen below:

For `haarcascade_frontalface_alt2` the human faces detection rates are:

- 100 for `human_files`
- 20 for `dog_files`

For `haarcascade_frontalface_default` the human faces detection rates are:

- 100 for `human_files`
- 58 for `dog_files`

For `haarcascade_frontalface_alt_tree` the human faces detection rates are:

- 51 for `human_files`
- 1 for `dog_files`

From all the results we see that probably the best method overall from the tested ones is based on the file `haarcascade_frontalface_alt` from the previous cells where the results showed high correctness of the human faces classifications and relatively low compared to other methods false alarm rates. Results for `haarcascade_frontalface_alt` were:

- 99 for `human_files`
- 11 for `dog_files`

The results should be probably improved as they could indeed frustrate the users, especially that a significant amount of dog images has actually human faces in them. Perhaps a good way to approach this problem would be also to try to search not only for human faces but actually for certain features of human faces such as eye, nose, mouth etc. and we know that these detectors are already available within OPENCV Haar files library. If human face correlates and coexists together with typical facial features then we can have a more robust detection for humans and a lesser false alarm rate for human faces detection in dog images.

In [8]:

```
## (Optional) TODO: Report the performance of another
## face detection algorithm on the LFW dataset
### Feel free to use as many code cells as needed.

human_files_short = human_files[:100]
dog_files_short = train_files[:100]

face_cascade_2 =
cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt2.xml')
face_cascade_default =
cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_default.xml')
face_cascade_alt_tree =
cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt_tree.xml')
```

```

def face_detector_2(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade_2.detectMultiScale(gray)
    return len(faces) > 0

countHumans = 0
countDogs = 0
for i in range(100):
    if face_detector_2(human_files_short[i]):
        countHumans += 1
    if face_detector_2(dog_files_short[i]):
        countDogs += 1

print('Alt2 Method: number of faces in human 100 images: {}'.format(countHumans))
print('Alt2 Method: number of faces in dog 100 images: {}'.format(countDogs))

def face_detector_default(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade_default.detectMultiScale(gray)
    return len(faces) > 0

countHumans = 0
countDogs = 0
for i in range(100):
    if face_detector_default(human_files_short[i]):
        countHumans += 1
    if face_detector_default(dog_files_short[i]):
        countDogs += 1

print('Default Method: number of faces in human 100 images: {}'.format(countHumans))
print('Default Method: number of faces in dog 100 images: {}'.format(countDogs))

def face_detector_alt_tree(img_path):
    img = cv2.imread(img_path)
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    faces = face_cascade_alt_tree.detectMultiScale(gray)
    return len(faces) > 0

countHumans = 0
countDogs = 0
for i in range(100):
    if face_detector_alt_tree(human_files_short[i]):
        countHumans += 1
    if face_detector_alt_tree(dog_files_short[i]):
        countDogs += 1

print('Alt_Tree Method: number of faces in human 100 images: {}'.format(countHumans))
print('Alt_Tree Method: number of faces in dog 100 images: {}'.format(countDogs))

```

Alt2 Method: number of faces in human 100 images: 100

Alt2 Method: number of faces in dog 100 images: 20

Default Method: number of faces in human 100 images: 100

Default Method: number of faces in dog 100 images: 20

Default Method: number of faces in dog 100 images: 58
Alt_Tree Method: number of faces in human 100 images: 51
Alt_Tree Method: number of faces in dog 100 images: 1

Step 2: Detect Dogs

In this section, we use a pre-trained [ResNet-50](#) model to detect dogs in images. Our first line of code downloads the ResNet-50 model, along with weights that have been trained on [ImageNet](#), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#). Given an image, this pre-trained ResNet-50 model returns a prediction (derived from the available categories in ImageNet) for the object that is contained in the image.

In [93]:

```
from keras.applications.resnet50 import ResNet50

# define ResNet50 model
ResNet50_model = ResNet50(weights='imagenet')
```

Pre-process the Data

When using TensorFlow as backend, Keras CNNs require a 4D array (which we'll also refer to as a 4D tensor) as input, with shape

\$\$ (\text{nb_samples}, \text{rows}, \text{columns}, \text{channels}), \$\$

where `nb_samples` corresponds to the total number of images (or samples), and `rows`, `columns`, and `channels` correspond to the number of rows, columns, and channels for each image, respectively.

The `path_to_tensor` function below takes a string-valued file path to a color image as input and returns a 4D tensor suitable for supplying to a Keras CNN. The function first loads the image and resizes it to a square image that is 224×224 pixels. Next, the image is converted to an array, which is then resized to a 4D tensor. In this case, since we are working with color images, each image has three channels. Likewise, since we are processing a single image (or sample), the returned tensor will always have shape

\$\$ (1, 224, 224, 3). \$\$

The `paths_to_tensor` function takes a numpy array of string-valued image paths as input and returns a 4D tensor with shape

\$\$ (\text{nb_samples}, 224, 224, 3). \$\$

Here, `nb_samples` is the number of samples, or number of images, in the supplied array of image paths. It is best to think of `nb_samples` as the number of 3D tensors (where each 3D tensor corresponds to a different image) in your dataset!

In [94]:

```
from keras.preprocessing import image
from tqdm import tqdm
```



```

def path_to_tensor(img_path):
    # loads RGB image as PIL.Image.Image type
    img = image.load_img(img_path, target_size=(224, 224))
    # convert PIL.Image.Image type to 3D tensor with shape (224, 224, 3)
    x = image.img_to_array(img)
    # convert 3D tensor to 4D tensor with shape (1, 224, 224, 3) and return
    4D tensor
    return np.expand_dims(x, axis=0)

def paths_to_tensor(img_paths):
    list_of_tensors = [path_to_tensor(img_path) for img_path in
    tqdm(img_paths)]
    return np.vstack(list_of_tensors)

```

Making Predictions with ResNet-50

Getting the 4D tensor ready for ResNet-50, and for any other pre-trained model in Keras, requires some additional processing. First, the RGB image is converted to BGR by reordering the channels. All pre-trained models have the additional normalization step that the mean pixel (expressed in RGB as $[103.939, 116.779, 123.68]$ and calculated from all pixels in all images in ImageNet) must be subtracted from every pixel in each image. This is implemented in the imported function `preprocess_input`. If you're curious, you can check the code for `preprocess_input` [here](#).

Now that we have a way to format our image for supplying to ResNet-50, we are now ready to use the model to extract the predictions. This is accomplished with the `predict` method, which returns an array whose i -th entry is the model's predicted probability that the image belongs to the i -th ImageNet category. This is implemented in the `ResNet50_predict_labels` function below.

By taking the argmax of the predicted probability vector, we obtain an integer corresponding to the model's predicted object class, which we can identify with an object category through the use of this [dictionary](#).

In [95]:

```

from keras.applications.resnet50 import preprocess_input,
decode_predictions

def ResNet50_predict_labels(img_path):
    # returns prediction vector for image located at img_path
    img = preprocess_input(path_to_tensor(img_path))
    return np.argmax(ResNet50_model.predict(img))

```

Write a Dog Detector

While looking at the [dictionary](#), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained ResNet-50 model, we need only check if the `ResNet50_predict_labels` function above returns a value between 151 and 268 (inclusive).

We use these ideas to complete the `dog_detector` function below, which returns `True` if a dog is detected in an image (and `False` if not).

In [96]:

```
### returns "True" if a dog is detected in the image stored at img_path
def dog_detector(img_path):
    prediction = ResNet50_predict_labels(img_path)
    return ((prediction <= 268) & (prediction >= 151))
```

(IMPLEMENTATION) Assess the Dog Detector

Question 3: Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

Answer: There is 0 of detected dogs in the `human_files_short` and there are 100 of detected dogs in the `dog_files_short`. This indicates an extremely robust dog detector function with all classifications for dogs correct and 0% of false alarm rate.

In [97]:

```
### TODO: Test the performance of the dog_detector function
### on the images in human_files_short and dog_files_short.

human_files_short = human_files[:100]
dog_files_short = train_files[:100]

countHumans = 0
countDogs = 0
for i in range(100):
    if dog_detector(human_files_short[i]):
        countHumans += 1
    if dog_detector(dog_files_short[i]):
        countDogs += 1

print('Dog Detector: number of dogs in 100 human images: {}'.format(countHumans))
print('Dog Detector: number of dogs in 100 dog images: {}'.format(countDogs))
```

```
Dog Detector: number of dogs in 100 human images: 0
Dog Detector: number of dogs in 100 dog images: 100
```

Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning yet!), and you must attain a test accuracy of at least 1%. In Step 5 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

Be careful with adding too many trainable layers! More parameters means longer training, which means you are more likely to need a GPU to accelerate the training process. Thankfully, Keras provides a handy estimate of the time that each epoch is likely to take; you can extrapolate this

estimate to figure out how long it will take for your algorithm to train.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have great difficulty in distinguishing between a Brittany and a Welsh Springer Spaniel.

Brittany	Welsh Springer Spaniel

It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

Curly-Coated Retriever	American Water Spaniel

Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

Yellow Labrador	Chocolate Labrador	Black Labrador

We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

Pre-process the Data

We rescale the images by dividing every pixel in every image by 255.

In [14]:

```
from PIL import ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True

# pre-process the data for Keras
train_tensors = paths_to_tensor(train_files).astype('float32')/255
valid_tensors = paths_to_tensor(valid_files).astype('float32')/255
test_tensors = paths_to_tensor(test_files).astype('float32')/255
```

```
100%|██████████| 6680/6680 [00:53<00:00, 125.74it/s]
100%|██████████| 835/835 [00:05<00:00, 140.86it/s]
100%|██████████| 836/836 [00:05<00:00, 141.77it/s]
```

(IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
model.summary()
```

We have imported some Python modules to get you started, but feel free to import as many modules

as you need. If you end up getting stuck, here's a hint that specifies a model that trains relatively fast on CPU and attains >1% test accuracy in 5 epochs:



Question 4: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. If you chose to use the hinted architecture above, describe why you think that CNN architecture should work well for the image classification task.

Answer: I decided to use the architecture which was hinted above. The convolution layers learn more and more advanced features of images as they propagate through the network. Max pooling layers help with overfitting reduction as they provide only an approximation passing to next layers of CNNs. The global average pooling takes the average value of each feature map and finally dense layer together with softmax activation outputs the prediction category of the input image.

In [15]:

```
from keras.layers import Conv2D, MaxPooling2D, GlobalAveragePooling2D
from keras.layers import Dropout, Flatten, Dense
from keras.models import Sequential

model = Sequential()

### TODO: Define your architecture.

#model.add(Conv2D(filters=16, kernel_size=2, padding='same', activation='relu',
#                input_shape=(223, 223, 3)))
model.add(Conv2D(filters=16, kernel_size=2, padding='same',
                activation='relu',
                input_shape=(224, 224, 3)))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters=32, kernel_size=2, padding='same',
                activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(Conv2D(filters=64, kernel_size=2, padding='same',
                activation='relu'))
model.add(MaxPooling2D(pool_size=2))
model.add(GlobalAveragePooling2D(data_format='channels_last'))
model.add(Dense(133, activation='softmax'))

model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 224, 224, 16)	208
max_pooling2d_2 (MaxPooling2D)	(None, 112, 112, 16)	0
conv2d_2 (Conv2D)	(None, 112, 112, 32)	2080
max_pooling2d_3 (MaxPooling2D)	(None, 56, 56, 32)	0
conv2d_3 (Conv2D)	(None, 56, 56, 64)	8256
max_pooling2d_4 (MaxPooling2D)	(None, 28, 28, 64)	0
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 64)	0

dense_1 (Dense)	(None, 133)	8645
=====		
Total params: 19,189		
Trainable params: 19,189		
Non-trainable params: 0		

Compile the Model

In [16]:

```
model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
```

(IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to [augment the training data](#), but this is not a requirement.

In [17]:

```
from keras.callbacks import ModelCheckpoint

### TODO: specify the number of epochs that you would like to use to train the model.

#epochs = ...
epochs = 10 #using 10 epochs shows no overfitting and the validation accuracy still keeps increasing.

### Do NOT modify the code below this line.

checkpointer =
ModelCheckpoint(filepath='saved_models/weights.best.from_scratch.hdf5',
                verbose=1, save_best_only=True)

model.fit(train_tensors, train_targets,
          validation_data=(valid_tensors, valid_targets),
          epochs=epochs, batch_size=20, callbacks=[checkpointer], verbose=1)
```

Train on 6680 samples, validate on 835 samples

Epoch 1/10

```
6660/6680 [=====>.] - ETA: 0s - loss: 4.8851 - acc: 0.0080
Epoch 00000: val_loss improved from inf to 4.86863, saving model to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 185s - loss: 4.8851 - acc: 0.0079 - val_loss: 4.8686 - val_acc: 0.0108
```

Epoch 2/10

```
6660/6680 [=====>.] - ETA: 0s - loss: 4.8614 - acc: 0.0113
Epoch 00001: val_loss improved from 4.86863 to 4.84848, saving model to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 185s - loss: 4.8613 - acc: 0.0112 - val_loss: 4.8485 - val_acc: 0.0228
```

Epoch 3/10

```

Epoch 3/10
6660/6680 [=====>.] - ETA: 0s - loss: 4.8216 - acc:
0.0173Epoch 00002: val_loss improved from 4.84848 to 4.81096, saving model
to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 186s - loss: 4.8218 - acc: 0.0
172 - val_loss: 4.8110 - val_acc: 0.0156
Epoch 4/10
6660/6680 [=====>.] - ETA: 0s - loss: 4.7839 - acc:
0.0162Epoch 00003: val_loss improved from 4.81096 to 4.79628, saving model
to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 186s - loss: 4.7839 - acc: 0.0
162 - val_loss: 4.7963 - val_acc: 0.0168
Epoch 5/10
6660/6680 [=====>.] - ETA: 0s - loss: 4.7561 - acc:
0.0210Epoch 00004: val_loss improved from 4.79628 to 4.77809, saving model
to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 185s - loss: 4.7559 - acc: 0.0
211 - val_loss: 4.7781 - val_acc: 0.0216
Epoch 6/10
6660/6680 [=====>.] - ETA: 0s - loss: 4.7259 - acc:
0.0234Epoch 00005: val_loss improved from 4.77809 to 4.75662, saving model
to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 185s - loss: 4.7253 - acc: 0.0
237 - val_loss: 4.7566 - val_acc: 0.0287
Epoch 7/10
6660/6680 [=====>.] - ETA: 0s - loss: 4.7006 - acc:
0.0261Epoch 00006: val_loss improved from 4.75662 to 4.73417, saving model
to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 185s - loss: 4.7006 - acc: 0.0
260 - val_loss: 4.7342 - val_acc: 0.0263
Epoch 8/10
6660/6680 [=====>.] - ETA: 0s - loss: 4.6747 - acc:
0.0279Epoch 00007: val_loss improved from 4.73417 to 4.71937, saving model
to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 187s - loss: 4.6745 - acc: 0.0
278 - val_loss: 4.7194 - val_acc: 0.0311
Epoch 9/10
6660/6680 [=====>.] - ETA: 0s - loss: 4.6543 - acc:
0.0327Epoch 00008: val_loss improved from 4.71937 to 4.70309, saving model
to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 188s - loss: 4.6545 - acc: 0.0
328 - val_loss: 4.7031 - val_acc: 0.0311
Epoch 10/10
6660/6680 [=====>.] - ETA: 0s - loss: 4.6332 - acc:
0.0353Epoch 00009: val_loss improved from 4.70309 to 4.69954, saving model
to saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 185s - loss: 4.6328 - acc: 0.0
353 - val_loss: 4.6995 - val_acc: 0.0371

```

Out[17]:

```
<keras.callbacks.History at 0x7f7631159ac8>
```

Load the Model with the Best Validation Loss

In [18]:

```
model.load_weights('saved_models/weights.best.from_scratch.hdf5')
```

Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 1%.

In [19]:

```
# get index of predicted dog breed for each image in test set
dog_breed_predictions = [np.argmax(model.predict(np.expand_dims(tensor, axis=0))) for tensor in test_tensors]

# report test accuracy
test_accuracy = 100*np.sum(np.array(dog_breed_predictions)==np.argmax(test_targets, axis=1))/len(dog_breed_predictions)
print('Test accuracy: %.4f%%' % test_accuracy)
```

Test accuracy: 3.3493%

Step 4: Use a CNN to Classify Dog Breeds

To reduce training time without sacrificing accuracy, we show you how to train a CNN using transfer learning. In the following step, you will get a chance to use transfer learning to train your own CNN.

Obtain Bottleneck Features

In [20]:

```
bottleneck_features = np.load('bottleneck_features/DogVGG16Data.npz')
train_VGG16 = bottleneck_features['train']
valid_VGG16 = bottleneck_features['valid']
test_VGG16 = bottleneck_features['test']
```

Model Architecture

The model uses the the pre-trained VGG-16 model as a fixed feature extractor, where the last convolutional output of VGG-16 is fed as input to our model. We only add a global average pooling layer and a fully connected layer, where the latter contains one node for each dog category and is equipped with a softmax.

In [21]:

```
VGG16_model = Sequential()
VGG16_model.add(GlobalAveragePooling2D(input_shape=train_VGG16.shape[1:]))
VGG16_model.add(Dense(133, activation='softmax'))

VGG16_model.summary()
```

Layer (type)	Output Shape	Param #
=====		
global_average_pooling2d_2 ((None, 512)	0
=====		
dense_2 (Dense)	(None, 133)	68229
=====		

Total params: 68,229
Trainable params: 68,229
Non-trainable params: 0

Compile the Model

In [22]:

```
VGG16_model.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
```

Train the Model

In [23]:

```
checkpointer =  
ModelCheckpoint(filepath='saved_models/weights.best.VGG16.hdf5',  
                verbose=1, save_best_only=True)  
  
VGG16_model.fit(train_VGG16, train_targets,  
                validation_data=(valid_VGG16, valid_targets),  
                epochs=20, batch_size=20, callbacks=[checkpointer], verbose=1)
```

Train on 6680 samples, validate on 835 samples

Epoch 1/20

6400/6680 [=====>...] - ETA: 0s - loss: 12.7187 - acc: 0.1033
Epoch 00000: val_loss improved from inf to 11.07293, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 12.6402 - acc: 0.1073 - val_loss: 11.0729 - val_acc: 0.1988

Epoch 2/20

6380/6680 [=====>...] - ETA: 0s - loss: 10.4074 - acc: 0.2641
Epoch 00001: val_loss improved from 11.07293 to 10.32878, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 10.3888 - acc: 0.2654 - val_loss: 10.3288 - val_acc: 0.2683

Epoch 3/20

6420/6680 [=====>...] - ETA: 0s - loss: 9.8682 - acc: 0.3313
Epoch 00002: val_loss improved from 10.32878 to 10.09807, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 9.8721 - acc: 0.3313 - val_loss: 10.0981 - val_acc: 0.2910

Epoch 4/20

6420/6680 [=====>...] - ETA: 0s - loss: 9.5609 - acc: 0.3615
Epoch 00003: val_loss improved from 10.09807 to 9.92161, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 9.5326 - acc: 0.3626 - val_loss: 9.9216 - val_acc: 0.3042

Epoch 5/20

6420/6680 [=====>...] - ETA: 0s - loss: 9.4060 - acc: 0.3849
Epoch 00004: val_loss improved from 9.92161 to 9.83586, saving model to saved_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 1s - loss: 9.4077 - acc: 0.3843 - val_loss: 9.8359 - val_acc: 0.3126

Epoch 6/20

6420/6680 [=====>...] - ETA: 0s - loss: 9.1548 - acc: 0.4002
Epoch 00005: val loss improved from 9.83586 to 9.49024, saving model


```
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 9.1351 - acc: 0.400
6 - val_loss: 9.4902 - val_acc: 0.3473
Epoch 7/20
6400/6680 [=====>..] - ETA: 0s - loss: 8.8332 - acc:
0.4222Epoch 00006: val_loss did not improve
6680/6680 [=====] - 1s - loss: 8.8584 - acc: 0.421
0 - val_loss: 9.5601 - val_acc: 0.3353
Epoch 8/20
6400/6680 [=====>..] - ETA: 0s - loss: 8.7699 - acc:
0.4355Epoch 00007: val_loss improved from 9.49024 to 9.36183, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.7597 - acc: 0.436
2 - val_loss: 9.3618 - val_acc: 0.3485
Epoch 9/20
6420/6680 [=====>..] - ETA: 0s - loss: 8.6901 - acc:
0.4435Epoch 00008: val_loss improved from 9.36183 to 9.21579, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.6873 - acc: 0.444
0 - val_loss: 9.2158 - val_acc: 0.3677
Epoch 10/20
6400/6680 [=====>..] - ETA: 0s - loss: 8.6327 - acc:
0.4488Epoch 00009: val_loss improved from 9.21579 to 9.21220, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.6332 - acc: 0.448
1 - val_loss: 9.2122 - val_acc: 0.3581
Epoch 11/20
6420/6680 [=====>..] - ETA: 0s - loss: 8.5628 - acc:
0.4561Epoch 00010: val_loss improved from 9.21220 to 9.20586, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.5477 - acc: 0.456
6 - val_loss: 9.2059 - val_acc: 0.3653
Epoch 12/20
6400/6680 [=====>..] - ETA: 0s - loss: 8.4472 - acc:
0.4597Epoch 00011: val_loss improved from 9.20586 to 9.02050, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.4549 - acc: 0.459
1 - val_loss: 9.0205 - val_acc: 0.3689
Epoch 13/20
6380/6680 [=====>..] - ETA: 0s - loss: 8.3188 - acc:
0.4704Epoch 00012: val_loss did not improve
6680/6680 [=====] - 1s - loss: 8.3145 - acc: 0.470
4 - val_loss: 9.0606 - val_acc: 0.3629
Epoch 14/20
6380/6680 [=====>..] - ETA: 0s - loss: 8.1875 - acc:
0.4746Epoch 00013: val_loss improved from 9.02050 to 8.88427, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.1681 - acc: 0.475
7 - val_loss: 8.8843 - val_acc: 0.3784
Epoch 15/20
6400/6680 [=====>..] - ETA: 0s - loss: 7.9910 - acc:
0.4892Epoch 00014: val_loss improved from 8.88427 to 8.73741, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 8.0099 - acc: 0.487
3 - val_loss: 8.7374 - val_acc: 0.3832
Epoch 16/20
6400/6680 [=====>..] - ETA: 0s - loss: 7.8570 - acc:
0.4977Epoch 00015: val_loss improved from 8.73741 to 8.61677, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 7.8587 - acc: 0.497
```

```

3 - val_loss: 8.6168 - val_acc: 0.3772
Epoch 17/20
6380/6680 [=====>..] - ETA: 0s - loss: 7.7762 - acc:
0.5074Epoch 00016: val_loss improved from 8.61677 to 8.54529, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 7.7763 - acc: 0.507
5 - val_loss: 8.5453 - val_acc: 0.3952
Epoch 18/20
6420/6680 [=====>..] - ETA: 0s - loss: 7.7719 - acc:
0.5084Epoch 00017: val_loss did not improve
6680/6680 [=====] - 1s - loss: 7.7635 - acc: 0.509
3 - val_loss: 8.6347 - val_acc: 0.3892
Epoch 19/20
6420/6680 [=====>..] - ETA: 0s - loss: 7.7463 - acc:
0.5134Epoch 00018: val_loss improved from 8.54529 to 8.50060, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 1s - loss: 7.7512 - acc: 0.513
0 - val_loss: 8.5006 - val_acc: 0.3976
Epoch 20/20
6420/6680 [=====>..] - ETA: 0s - loss: 7.7291 - acc:
0.5150Epoch 00019: val_loss did not improve
6680/6680 [=====] - 1s - loss: 7.7399 - acc: 0.514
2 - val_loss: 8.5233 - val_acc: 0.3976

```

Out[23]:

```
<keras.callbacks.History at 0x7f7630738fd0>
```

Load the Model with the Best Validation Loss

In [24]:

```
VGG16_model.load_weights('saved_models/weights.best.VGG16.hdf5')
```

Test the Model

Now, we can use the CNN to test how well it identifies breed within our test dataset of dog images. We print the test accuracy below.

In [25]:

```

# get index of predicted dog breed for each image in test set
VGG16_predictions = [np.argmax(VGG16_model.predict(np.expand_dims(feature,
axis=0))) for feature in test_VGG16]

# report test accuracy
test_accuracy =
100*np.sum(np.array(VGG16_predictions)==np.argmax(test_targets, axis=1))/le
n(VGG16_predictions)
print('Test accuracy: %.4f%%' % test_accuracy)

```

Test accuracy: 41.8660%

Predict Dog Breed with the Model

In [26]:

```

from extract_bottleneck_features import *

def VGG16_predict_breed(img_path):
    # extract bottleneck features
    bottleneck_feature = extract_VGG16(path_to_tensor(img_path))
    # obtain predicted vector
    predicted_vector = VGG16_model.predict(bottleneck_feature)
    # return dog breed that is predicted by the model
    return dog_names[np.argmax(predicted_vector)]

```

Step 5: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

In Step 4, we used transfer learning to create a CNN using VGG-16 bottleneck features. In this section, you must use the bottleneck features from a different pre-trained model. To make things easier for you, we have pre-computed the features for all of the networks that are currently available in Keras:

- [VGG-19](#) bottleneck features
- [ResNet-50](#) bottleneck features
- [Inception](#) bottleneck features
- [Xception](#) bottleneck features

The files are encoded as such:

```
Dog{network}Data.npz
```

where {network}, in the above filename, can be one of VGG19, Resnet50, InceptionV3, or Xception. Pick one of the above architectures, download the corresponding bottleneck features, and store the downloaded file in the `bottleneck_features/` folder in the repository.

(IMPLEMENTATION) Obtain Bottleneck Features

In the code block below, extract the bottleneck features corresponding to the train, test, and validation sets by running the following:

```

bottleneck_features =
np.load('bottleneck_features/Dog{network}Data.npz')
train_{network} = bottleneck_features['train']
valid_{network} = bottleneck_features['valid']
test_{network} = bottleneck_features['test']

```

In [70]:

```

### TODO: Obtain bottleneck features from another pre-trained CNN.

bottleneck_features = np.load('bottleneck_features/DogResnet50Data.npz')
train_Resnet50K = bottleneck_features['train']

```

```
valid_Resnet50K = bottleneck_features['valid']
test_Resnet50K = bottleneck_features['test']
```

(IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
<your model's name>.summary()
```

Question 5: Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

Answer: I have selected the Resnet50 model as the base model for transfer learning. I have added to it a Global Average subsampling to the fully connected dense layer. This decreases the dimensionality by a large factor while still preserving the features. The final layer is dense and has 133 outputs which correspond to the dog breeds. I used for that the softmax activation function.

Resnet50 performs much better as previous models as it is one of the best models available as proven in the ImageNet competition. This model is much deeper, has many more layers and also to avoid the loss of weights which would have to propagate through so many layers there are some interconnections and some weights skip many layers for that. This skipping mechanism allowed for integration of so many layers and hence to store really deep features in the network. This is probably why it outperforms the previously tried approaches.

In [71]:

```
### TODO: Define your architecture.

Resnet50K_model = Sequential()
Resnet50K_model.add(GlobalAveragePooling2D(input_shape=train_Resnet50K.shape[1:]))
Resnet50K_model.add(Dense(133, activation='softmax'))

Resnet50K_model.summary()
```

Layer (type)	Output Shape	Param #
global_average_pooling2d_5 ((None, 2048)		0
dense_5 (Dense)	(None, 133)	272517

Total params: 272,517
Trainable params: 272,517
Non-trainable params: 0

(IMPLEMENTATION) Compile the Model

In [72]:

```
### TODO: Compile the model.

Resnet50K_model.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
```

```
resnet50_model.compile(loss=categorical_crossentropy, optimizer=ImpProp  
, metrics=['accuracy'])
```

(IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to [augment the training data](#), but this is not a requirement.

In [73]:

```
### TODO: Train the model.
```

```
checkpointer =  
ModelCheckpoint(filepath='saved_models/weights.best.Resnet50.hdf5',  
                verbose=1, save_best_only=True)  
  
Resnet50K_model.fit(train_Resnet50K, train_targets,  
                    validation_data=(valid_Resnet50K, valid_targets),  
                    epochs=20, batch_size=20, callbacks=[checkpointer], verbose=1)
```

Train on 6680 samples, validate on 835 samples

Epoch 1/20

```
6580/6680 [=====>.] - ETA: 0s - loss: 1.6444 - acc:  
0.5954Epoch 00000: val_loss improved from inf to 0.81384, saving model to s  
aved_models/weights.best.Resnet50.hdf5
```

```
6680/6680 [=====] - 9s - loss: 1.6306 - acc: 0.598  
8 - val_loss: 0.8138 - val_acc: 0.7353
```

Epoch 2/20

```
6580/6680 [=====>.] - ETA: 0s - loss: 0.4438 - acc:  
0.8603Epoch 00001: val_loss improved from 0.81384 to 0.62251, saving model  
to saved_models/weights.best.Resnet50.hdf5
```

```
6680/6680 [=====] - 2s - loss: 0.4428 - acc: 0.861  
2 - val_loss: 0.6225 - val_acc: 0.7988
```

Epoch 3/20

```
6580/6680 [=====>.] - ETA: 0s - loss: 0.2593 - acc:  
0.9229Epoch 00002: val_loss did not improve
```

```
6680/6680 [=====] - 2s - loss: 0.2611 - acc: 0.922  
6 - val_loss: 0.6895 - val_acc: 0.8060
```

Epoch 4/20

```
6580/6680 [=====>.] - ETA: 0s - loss: 0.1805 - acc:  
0.9459Epoch 00003: val_loss did not improve
```

```
6680/6680 [=====] - 2s - loss: 0.1803 - acc: 0.946  
0 - val_loss: 0.6361 - val_acc: 0.8180
```

Epoch 5/20

```
6580/6680 [=====>.] - ETA: 0s - loss: 0.1234 - acc:  
0.9616Epoch 00004: val_loss improved from 0.62251 to 0.60875, saving model  
to saved_models/weights.best.Resnet50.hdf5
```

```
6680/6680 [=====] - 2s - loss: 0.1235 - acc: 0.961  
5 - val_loss: 0.6087 - val_acc: 0.8287
```

Epoch 6/20

```
6580/6680 [=====>.] - ETA: 0s - loss: 0.0885 - acc:  
0.9743Epoch 00005: val_loss did not improve
```

```
6680/6680 [=====] - 2s - loss: 0.0879 - acc: 0.974  
4 - val_loss: 0.6940 - val_acc: 0.8156
```

Epoch 7/20

```
6580/6680 [=====>.] - ETA: 0s - loss: 0.0643 - acc:  
0.9807Epoch 00006: val_loss did not improve
```

```
6680/6680 [=====] - 2s - loss: 0.0628 - acc: 0.980
```

```
6680/6680 [=====] - 2s - loss: 0.0639 - acc: 0.980
8 - val_loss: 0.6798 - val_acc: 0.8120
Epoch 8/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0469 - acc:
0.9854Epoch 00007: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0474 - acc: 0.985
2 - val_loss: 0.6690 - val_acc: 0.8323
Epoch 9/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0352 - acc:
0.9901Epoch 00008: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0354 - acc: 0.990
1 - val_loss: 0.7124 - val_acc: 0.8275
Epoch 10/20
6560/6680 [=====>.] - ETA: 0s - loss: 0.0255 - acc:
0.9936Epoch 00009: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0258 - acc: 0.993
6 - val_loss: 0.7806 - val_acc: 0.8299
Epoch 11/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0230 - acc:
0.9938Epoch 00010: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0229 - acc: 0.993
9 - val_loss: 0.7172 - val_acc: 0.8371
Epoch 12/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0190 - acc:
0.9947Epoch 00011: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0188 - acc: 0.994
8 - val_loss: 0.7725 - val_acc: 0.8347
Epoch 13/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0136 - acc:
0.9973Epoch 00012: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0135 - acc: 0.997
3 - val_loss: 0.8033 - val_acc: 0.8287
Epoch 14/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0119 - acc:
0.9976Epoch 00013: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0120 - acc: 0.997
5 - val_loss: 0.8288 - val_acc: 0.8263
Epoch 15/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0093 - acc:
0.9974Epoch 00014: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0093 - acc: 0.997
5 - val_loss: 0.8492 - val_acc: 0.8240
Epoch 16/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0091 - acc:
0.9980Epoch 00015: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0094 - acc: 0.997
9 - val_loss: 0.8313 - val_acc: 0.8180
Epoch 17/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0071 - acc:
0.9982Epoch 00016: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0071 - acc: 0.998
1 - val_loss: 0.8537 - val_acc: 0.8359
Epoch 18/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0064 - acc:
0.9985Epoch 00017: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0066 - acc: 0.998
4 - val_loss: 0.8226 - val_acc: 0.8395
Epoch 19/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0067 - acc:
0.9983Epoch 00018: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0066 - acc: 0.998
```

```

0000/0000 [-----] 2s - loss: 0.0000 - acc: 0.998
4 - val_loss: 0.8520 - val_acc: 0.8299
Epoch 20/20
6580/6680 [=====>.] - ETA: 0s - loss: 0.0051 - acc:
0.9983Epoch 00019: val_loss did not improve
6680/6680 [=====] - 2s - loss: 0.0051 - acc: 0.998
4 - val_loss: 0.9065 - val_acc: 0.8240

```

Out[73]:

```
<keras.callbacks.History at 0x7f739cf70f98>
```

(IMPLEMENTATION) Load the Model with the Best Validation Loss

In [74]:

```

### TODO: Load the model weights with the best validation loss.

Resnet50K_model.load_weights('saved_models/weights.best.Resnet50.hdf5')

```

(IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 60%.

In [75]:

```

### TODO: Calculate classification accuracy on the test dataset.

# get index of predicted dog breed for each image in test set
Resnet50K_predictions = [np.argmax(Resnet50K_model.predict(np.expand_dims(f
eature, axis=0))) for feature in test_Resnet50K]

# report test accuracy
test_accuracy = 100*np.sum(np.array(Resnet50K_predictions)==np.argmax(test_
targets, axis=1))/len(Resnet50K_predictions)
print('Test accuracy: %.4f%%' % test_accuracy)

```

```
Test accuracy: 82.0574%
```

(IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed (Affenpinscher, Afghan_hound, etc) that is predicted by your model.

Similar to the analogous function in Step 5, your function should have three steps:

1. Extract the bottleneck features corresponding to the chosen CNN model.
2. Supply the bottleneck features as input to the model to return the predicted vector. Note that the argmax of this prediction vector gives the index of the predicted dog breed.
3. Use the `dog_names` array defined in Step 0 of this notebook to return the corresponding breed.

The functions to extract the bottleneck features can be found in `extract_bottleneck_features.py`, and they have been imported in an earlier code cell. To obtain the bottleneck features corresponding to your chosen CNN architecture, you need to use the

function

```
extract_{network}
```

where {network}, in the above filename, should be one of VGG19, Resnet50, InceptionV3, or Xception.

In [101]:

```
### TODO: Write a function that takes a path to an image as input
### and returns the dog breed that is predicted by the model.

from extract_bottleneck_features import *

def Resnet50K_breed(img_path):
    # extract bottleneck features
    bottleneck_feature = extract_Resnet50(path_to_tensor(img_path))
    # obtain predicted vector
    predicted_vector = Resnet50K_model.predict(bottleneck_feature)
    # return dog breed that is predicted by the model
    return dog_names[np.argmax(predicted_vector)],
np.argmax(predicted_vector)
```

Step 6: Write your Algorithm

Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then,

- if a **dog** is detected in the image, return the predicted breed.
- if a **human** is detected in the image, return the resembling dog breed.
- if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `dog_detector` functions developed above. You are **required** to use your CNN from Step 5 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!



(IMPLEMENTATION) Write your Algorithm

In [105]:

```
### TODO: Write your algorithm.
### Feel free to use as many code cells as needed.

from IPython.display import Image, display
import matplotlib.image as mpimg

def dog_app(img_path):
```



```

if face_detector(img_path):
    breed, breed_id = Resnet50K_breed(img_path) #returns breed
    message = 'Hello Human!\nYou look like a ...'\n' + breed
elif dog_detector(img_path):
    breed, breed_id = Resnet50K_breed(img_path) #returns breed
    message = 'Hello Dog!\nYou look like a ...'\n' + breed
else:
    breed = 'Error: cannot identify the breed'
    breed_id = -1
    message = 'Hello!\n' + breed

return message, breed, breed_id

```

Step 7: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that **you** look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

(IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

Question 6: Is the output better than you expected :) ? Or worse :(? Provide at least three possible points of improvement for your algorithm.

Answer: Algorithms works pretty well for distinguishing between humans, dogs and neither of those categories. Additionally it performs reasonably well for dogs breeds identification. It is however difficult to judge whether the performance is satisfactory when it comes to matching a dog breed for a human face. Perhaps it works okay but for me it is very difficult to judge this as I see very little similarity.

Possible improvements:

- Augmenting the training dataset by adding versions of the same images with different rotations, different positions in the image or different sizes. This would add a higher degree of generalization to our algorithms.
- Training the whole model instead of just finetuning it.
- Playing with different network architectures, activation functions etc. might bring better results.

In [107]:

```

## TODO: Execute your algorithm from Step 6 on
## at least 6 images on your computer.
## Feel free to use as many code cells as needed.

test_files = np.array(glob("test_images/*"))

def get_sample(breed_id):
    i = 0
    for bid in train_targets:
        if np.argmax(bid) == breed_id:

```

```

    if np.argmax(scores) == breed_id:
        break
    i += 1
    return train_files[i]

for img_path in test_files:




    message, breed, breed_id = dog_app(img_path)

    fig = plt.figure()
    a = fig.add_subplot(1,2,1)
    img = mpimg.imread(img_path)
    imgplot = plt.imshow(img)
    a.set_title(message)


    if breed_id != -1:
        a = fig.add_subplot(1,2,2)
        sample_path = get_sample(breed_id)
        sample_img = mpimg.imread(sample_path)
        imgplot = plt.imshow(sample_img)
        a.set_title("Sample for {}".format(breed))

plt.show()


```




Hello Human!
You look like a ...
American_foxhound




Sample for American_foxhound




Hello Dog!
You look like a ...
Entlebucher_mountain_dog




Hello Dog!
You look like a ...
English_setter



Sample for English_setter



Hello Dog!
You look like a ...
Bernese_mountain_dog



Sample for Bernese_mountain_dog




Hello Dog!
You look like a ...
Chesapeake_bay_retriever



Sample for Chesapeake_bay_retriever





Hello Human!
You look like a ...
Silky_terrier

